Baifeng Shi

bfshi@pku.edu.cn | (+86) 18811630350
No.5 Yiheyuan Rd., Beijing, China, 100871
https://bfshi.github.io/

EDUCATION BACKGROUND

Peking University, Beijing, China 09/2017 - 06/2021 (expected)
B.S., Computer Science

- Overall GPA:**3.75** / 4
- Standard Tests: TOEFL 108 (Speaking 23) / GRE 332 (AW 3.5)
- Programming languages: C/C++, Python, MATLAB, R, LaTeX
- Awards and Honors:

  Gold Medal (3 / 360), Chinese Physics Olympiad final contest 10/2016
  First Class Golden Award, 12th Pan-Pearl River Delta and Chinese Elite Schools Physics Olympiad 02/2016
  EECS Dean Scholarship, Peking University 09/2017
  Merit Student, Peking University 09/2018 & 09/2020

RESEARCH APPOINTMENTS

- **University of California, Berkeley**, Research Intern 03/2020 - 11/2020
  Advisor: Dr. Huijuan Xu & Prof. Trevor Darrell
- **Microsoft Research Asia**, Research Intern 09/2019 - 12/2020
  Advisor: Dr. Qi Dai & Dr. Jingdong Wang
- **Peking University**, Research Intern 09/2018 - 09/2019
  Advisor: Prof. Yadong Mu

PUBLICATIONS

- **Baifeng Shi**, Judy Hoffman, Kate Saenko, Trevor Darrell, Huijuan Xu, *Auxiliary Task Reweighting for Minimum-data Learning*, **NeurIPS 2020**
- Zhekun Luo, Devin Guillory, **Baifeng Shi**, Wei Ke, Fang Wan, Trevor Darrell, Huijuan Xu, *Weakly-Supervised Action Localization with Expectation-Maximization Multi-Instance Learning*, **ECCV 2020**
- **Baifeng Shi***, Dinghuai Zhang*, Qi Dai, Zhanxing Zhu, Yadong Mu, Jingdong Wang, *Informative Dropout for Robust Representation Learning: A Shape-bias Perspective*, **ICML 2020**
- **Baifeng Shi**, Qi Dai, Jingdong Wang, Yadong Mu, *Weakly-Supervised Action Localization by Generative Attention Modeling*, **CVPR 2020**
- **Baifeng Shi**, Qi Dai, Judy Hoffman, Kate Saenko, Trevor Darrell, Huijuan Xu, *Temporal Action Detection with Multi-level Supervision*, **submitted to CVPR 2021**

RESEARCH EXPERIENCE

Independent research, advised by Dr. Huijuan Xu and Prof. Trevor Darrell *UC Berkeley*
Unsupervised Foreground Mining for Omni-supervised Action Localization 06/2020 – 11/2020

- Built the first baselines for semi-supervised and omni-supervised action localization.
- Designed error analysis to find the main sources of error in the baseline models.
- Proposed to solve the action incompleteness problem in the semi-supervised baseline by learning object-centric representations. Built a structural causal model of the foreground/background action in neighboring frames, and proposed to detect foreground objects by minimizing the conditional mutual information between foreground and background motion.
- Proposed to solve the action-context confusion problem in the omni-supervised baseline by designing an information bottleneck to discard scene information while preserve action information.
- One paper is submitted to CVPR 2021.

Independent research, advised by Dr. Huijuan Xu and Prof. Trevor Darrell *UC Berkeley*
Auxiliary Task Reweighting for Minimum-data Learning 03/2020 – 06/2020

- Addressed the problem of automatically reweighting multiple auxiliary tasks to learn the main task with minimum information (supervision).
- Exploited the key insight that *information required for inference can be reduced by a good prior*, and formulated the problem as optimizing the KL divergence between the true prior and the surrogate prior given by the weighted likelihood of auxiliary tasks.
- Utilized tools and concepts including Fisher score and Langevin dynamics, and further simplified the optimization of the KL divergence into minimizing the l2 distance between main/auxiliary task gradients, which gives a light-weight algorithm to reweight auxiliary tasks on-the-fly.
- Derived theoretical guarantees that our algorithm finds the optimal task weights up to a small error.

- Experimentally observed that our algorithm finds the optimal task weights and minimizes the data requirement under various settings, *e.g.*, semi-supervised learning, domain generalization, and multi-label classification.
- The work is summarized in a paper and accepted to NeurIPS 2020.

Independent research, mentored by Dr. Qi Dai, Dr. Jingdong Wang, and Prof. Yadong Mu *Microsoft Research Asia*
Human Vision Inspired Shape-bias for Model Robustness 11/2019 – 02/2020

- Analyzed the relationship between the texture-bias of CNN and its multiple kinds of vulnerability.
- Proposed to discriminate texture from shape by the self-information in an image, resembling the mechanism of saliency detection and eye movement in the human visual system.
- Proposed a Dropout-like algorithm to de-correlate the model output with the texture information in the input, thus enhancing the shape-bias of the model.
- Conducted experiments under different scenarios (domain generalization, few-shot classification, image corruption, and adversarial perturbation) and observed a universal improvement in model robustness.
- The work is summarized in a paper and accepted to ICML 2020.

Independent research, mentored by Dr. Qi Dai, Dr. Jingdong Wang, and Prof. Yadong Mu *Microsoft Research Asia*
Weakly Supervised Action Localization 08/2019 – 11/2019

- Tackled with a common challenge in weakly supervised action localization, namely action-context confusion.
- Built a probabilistic graphical model and formulated the problem as modeling the frame-wise class-agnostic likelihood and optimizing the maximum a posteriori (MAP) estimation.
- Proposed to separate foreground and context by modeling the appearance-level frame likelihood using a generative model, *viz.* conditional variational auto-encoder (CVAE).
- Improved the results on two common datasets by a large margin (10% relative improvement).
- The work is summarized in a paper and accepted to CVPR 2020.

Independent research, advised by Prof. Yadong Mu *Peking University*
Fast Video Understanding with Reinforcement Learning 10/2018 – 08/2019

- Proposed an algorithm to use as few frames as possible to classify a video while preserving an expected accuracy.
- Formulated the reward function as the Lagrangian of the constrained optimization problem, and optimized it using reinforcement learning algorithm, *viz.* soft actor-critic (SAC).
- Boost the original algorithm by 400% while preserving the accuracy.

SELECTED PROJECTS

Neural Decoding of Rhesus Monkey's Primary Motor Cortex (M1)
Advised by Prof. Bo Hong

- Preprocessed the spike rate and kinematics record from raw data.
- Decoded hand movement from the population code of primary motor cortex using LASSO.
- Estimated the tuning curve of M1 neurons and their preferred direction.
- Took the temporal consistency of motion into consideration, and use Kalman Filter to improve the decoding accuracy.

An Upper Bound on Adversarial Robustness and its Dependency on Model's Performance
Advised by Prof. Liwei Wang and Prof. John Hopcroft

- Theoretically proved an upper bound of adversarial robustness of an arbitrary model and analyzed its dependency on model's performance.
- Showed as the main result that this upper bound will get lower as performance improves.
- Conducted several experiments and heuristically proved our claim.

Towards Efficient methods for Margin Maximization in ReLU Networks
Advised by Prof. Liwei Wang

- Developed an algorithm to improve adversarial robustness of ReLU networks based on spectral norm regularization.
- Designed an algorithm to calculate spectral norm of convolutional layers more accurately and efficiently.